



CodeGeeX: A Pre-Trained Model for Code Generation with Multilingual Benchmarking on HumanEval-X

Qinkai Zheng*

qinkai@tsinghua.edu.cn
Tsinghua University
Beijing, China

Yuxiao Dong[†]

yuxiaod@tsinghua.edu.cn
Tsinghua University
Beijing, China

Lei Shen

lei.shen@zhipuai.cn
Zhipu.AI
Beijing, China

Yang Li

yang.li@zhipuai.cn
Zhipu.AI
Beijing, China

Xiao Xia*

xiax19@mails.tsinghua.edu.cn
Tsinghua University
Beijing, China

Shan Wang

shan.wang@zhipuai.cn
Zhipu.AI
Beijing, China

Zihan Wang

zhwang19@mails.tsinghua.edu.cn
Tsinghua University
Beijing, China

Teng Su

suteng@huawei.com
Huawei
Hangzhou, China

Jie Tang^{†‡}

jietang@tsinghua.edu.cn
Tsinghua University
Beijing, China

Xu Zou

zoux18@mails.tsinghua.edu.cn
Tsinghua University
Beijing, China

Yufei Xue

yufei.xue@zhipuai.cn
Zhipu.AI
Beijing, China

Andi Wang

andi.wang@zhipuai.cn
Zhipu.AI
Beijing, China

Zhilin Yang[†]

zhiliny@tsinghua.edu.cn
Tsinghua University
Beijing, China

ABSTRACT

Large pre-trained code generation models, such as OpenAI Codex, can generate syntax- and function-correct code, making the coding of programmers more productive. In this paper, we introduce CodeGeeX, a multilingual model with 13 billion parameters for code generation. CodeGeeX is pre-trained on 850 billion tokens of 23 programming languages as of June 2022. Our extensive experiments suggest that CodeGeeX outperforms multilingual code models of similar scale for both the tasks of code generation and translation on HumanEval-X. Building upon HumanEval (Python only), we develop the HumanEval-X benchmark for evaluating multilingual models by hand-writing the solutions in C++, Java, JavaScript, and Go. In addition, we build CodeGeeX-based extensions on Visual Studio Code, JetBrains, and Cloud Studio, generating 8 billion tokens for tens of thousands of active users per week. Our user study demonstrates that CodeGeeX can help to increase coding efficiency

*QZ and XX contributed equally.

[†]Team Leads: YD, ZY, and JT.

[‡]Corresponding Author: JT.



This work is licensed under a Creative Commons Attribution International 4.0 License.

for 83.4% of its users. Finally, CodeGeeX is publicly accessible since Sep. 2022, we open-sourced its code, model weights, API, extensions, and HumanEval-X at <https://github.com/THUDM/CodeGeeX>.

KEYWORDS

code generation, pre-trained model, large language model

ACM Reference Format:

Qinkai Zheng, Xiao Xia, Xu Zou, Yuxiao Dong, Shan Wang, Yufei Xue, Lei Shen, Zihan Wang, Andi Wang, Yang Li, Teng Su, Zhilin Yang, and Jie Tang. 2023. CodeGeeX: A Pre-Trained Model for Code Generation with Multilingual Benchmarking on HumanEval-X. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD '23)*, August 6–10, 2023, Long Beach, CA, USA. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3580305.3599790>

1 INTRODUCTION

Given the description such as “write a factorial function”, can the machine automatically generate an executable program that addresses this need? This problem has been explored since the early days of computer science in the 1960s [31, 37]. From LISP-based pioneering deductive synthesis approaches [31, 37] to modern program synthesis systems [23, 30], to end-to-end code generation via deep neural networks [19, 32, 33], tremendous efforts have been made to enable machines to automatically write correct programs as part of the quest to artificial general intelligence.

By treating programs as language sequences, neural sequential architectures, such as recurrent neural networks and transformer [36], can be naturally applied to code generation. Notably, the OpenAI Codex [7] model (Python only) with 12 billion (12B) parameters pioneered and demonstrated the potential of large code generation models pre-trained on billions of lines of public code. By using the generative pre-training strategy, Codex can solve introductory-level programming problems in Python with a high probability. Research studies [46] also show that 88% of users of GitHub Copilot—a paid service powered by Codex—feel more productive when coding with it. Since then, large pre-trained code models have been extensively developed, including DeepMind AlphaCode [17], Salesforce CodeGen [20], Meta InCoder [10], and Google PaLM-Coder-540B [8].

In this work, we present CodeGeeX, a multilingual code generation model with 13 billion parameters, pre-trained on a large code corpus of 23 programming languages. It was trained on more than 850 billion tokens on a cluster of 1,536 Ascend 910 AI Processors between April and June 2022, and was publicly released in Sep. 2022 (Cf. the GitHub repo). CodeGeeX has the following properties. First, different from Codex in [7], both CodeGeeX—the model itself—and how such scale of code models can be pre-trained are open-sourced, facilitating the understanding and advances in pre-trained code models. Second, in addition to code generation and code completion as Codex and others, CodeGeeX supports the tasks of code explanation and code translation between language pairs (Cf. Figure 1 (a)). Third, it offers consistent performance advantages over well-known *multilingual* code generation models of the same scale, including CodeGen-16B, GPT-NeoX-20B, InCoder-6.7B, and GPT-J-6B (Cf. Figure 1 (b) and (c)).

We build the free CodeGeeX extension in several IDEs, currently including Visual Studio Code, JetBrains, and Tencent Cloud Studio (a Web IDE). It supports several different modes—code completion, function-level generation, code translation, code explanation, and customizable prompting—to help users’ programming tasks in real-time. Since its release, there are tens of thousands of daily active users, each of which on average makes 200+ API calls per weekday. As of this writing, the CodeGeeX model generates 8 billion tokens per week. Our user survey suggests that 83.4% of users feel the CodeGeeX extensions improve their programming efficiency.

Additionally, we develop the HumanEval-X benchmark for evaluating multilingual code models as 1) HumanEval [7]—developed by OpenAI for evaluating Codex—and other benchmarks [2, 13, 20] only consist of programming problems in a single language and 2) existing multilingual datasets [18, 28, 45] use string similarity metrics like BLEU [21] for evaluation rather than really verify the functional correctness of generated code. Specifically, for each problem—defined only for Python—in HumanEval, we manually rewrite its prompt, canonical solution, and test cases in C++, Java, JavaScript, and Go. In total, HumanEval-X covers 820 hand-written problem-solution pairs (164 problems, each of which has solutions in 5 languages). Importantly, HumanEval-X supports the evaluation of both code generation and translation between different languages. Our contributions can be summarized as follows:

- We develop and release CodeGeeX, a 13B pre-trained 23-language code generation model that demonstrates consistent outperformance on code generation and translation over its multilingual baselines of the same scale.
- We build the CodeGeeX extensions on VS Code, JetBrains, and Tencent Cloud Studio. Compared to Copilot, it supports more diverse functions, including code completion, generation, translation, and explanation. According to the user survey, CodeGeeX can improve the coding efficiency for 83.4% of its users.
- We hand-craft the HumanEval-X benchmark to evaluate multilingual code models for the tasks of code generation and translation in terms of functional correctness, facilitating the understanding and development of pre-trained (multilingual) code models.

2 THE CODEGEEX MODEL

CodeGeeX is a multilingual code generation model with 13 billion (13B) parameters, pre-trained on a large code corpus of 23 programming languages. As of June 22, 2022, CodeGeeX has been trained on more than 850 billion tokens on a cluster of 1,536 Ascend 910 AI Processors for over two months. We introduce the CodeGeeX model and its design choices. The consensus reality is that it is computationally unaffordable to test different architectural designs for large pre-trained models [6, 8, 42, 44].

2.1 CodeGeeX’s Architecture

Transformer Backbone. Similar to recent pre-trained models, such as GPT-3 [6], PaLM [8], and Codex [7], CodeGeeX follows the generative pre-trained transformer (GPT) [25] with the decoder-only style for autoregressive (programming) language modeling. The core architecture of CodeGeeX is a 39-layer transformer decoder. In each transformer layer (in Figure 2), we apply a multi-head self-attention mechanism [36] followed by MLP layers, together with layer normalization [3] and residual connection [12]. We use an approximation of GELU (Gaussian Linear Units) operation [14], namely FastGELU, which is more efficient under Ascend 910:

$$\text{FastGELU}(X_i) = \frac{X_i}{1 + \exp(-1.702 * |X_i|) * \exp(0.851 * (X_i - |X_i|))} \quad (1)$$

Generative Pre-Training Objective. By adopting the GPT paradigm [7, 26], we train the model on a large amount of unlabeled code data. The principle is to iteratively take code tokens as input, predict the next token, and compare it with the ground truth. Specifically, for any input sequence $\{x_1, x_2, \dots, x_n\}$ of length n , the output of CodeGeeX is a probability distribution of the next token $\mathbb{P}(x_{n+1}|x_1, x_2, \dots, x_n, \Theta) = p_{n+1} \in [0, 1]^{1 \times v}$, where Θ represents all parameters of the model and v is the vocabulary size. By comparing it with the ground-truth tokens, i.e., one-hot vector $y_{n+1} \in \{0, 1\}^{1 \times v}$, we can optimize the cumulative cross-entropy loss:

$$\mathcal{L} = - \sum_{n=1}^{N-1} y_{n+1} \log \mathbb{P}(x_{n+1}|x_1, x_2, \dots, x_n, \Theta) \quad (2)$$

Top Query Layer and Decoding. The original GPT model uses a pooler function to obtain the final output. We use an extra query layer [43] on top of all other transformer layers to obtain the final embedding through attention. As shown in Figure 2, the input of the top query layer replaces the query input X_{in} by the query

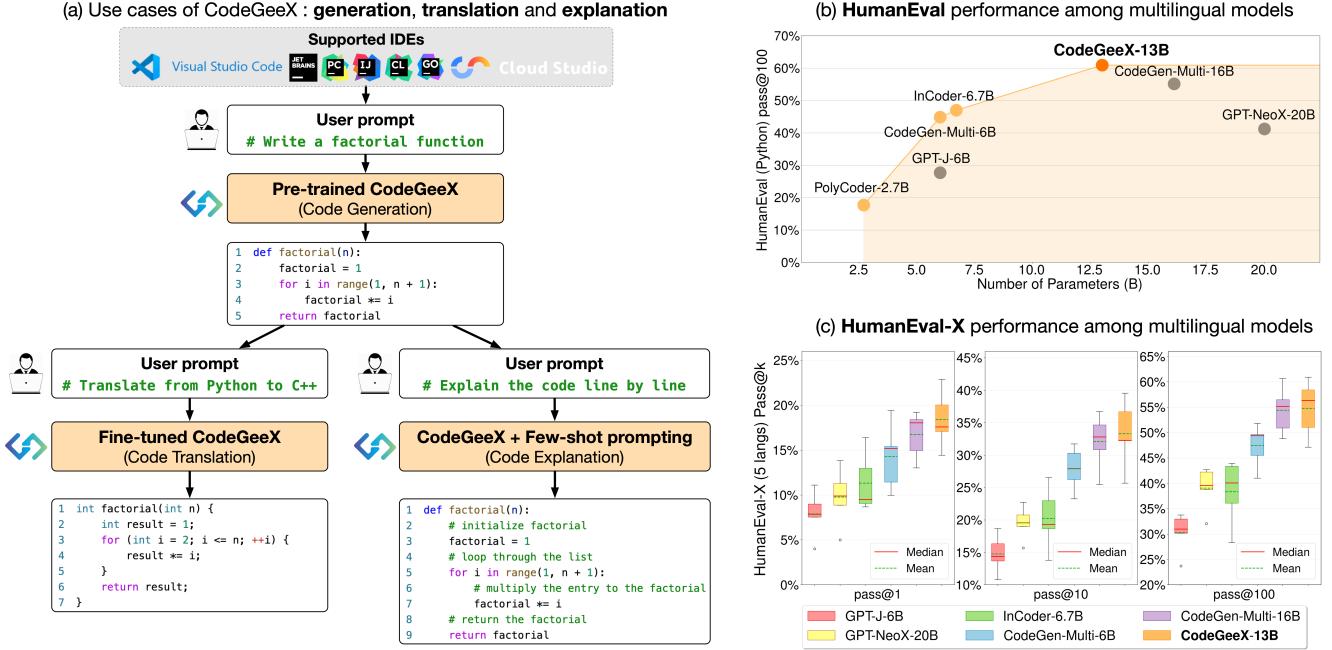


Figure 1: Summary of CodeGeeX. (a): In supported IDEs, users can interact with CodeGeeX by providing prompts. Different models are used to support three tasks: code generation, code translation and code explanation. (b) and (c): In HumanEval and our newly proposed HumanEval-X, CodeGeeX shows promising multilingual ability and consistently outperforms other multilingual code generation models.

Table 1: Large pre-trained language models related to programming languages in the literature.

	Model Properties			Dataset			Evaluation		
	Open	Multi-lingual	# Params	Source	Languages	Size	Multilingual Evaluation	Translation	Benchmark
Codex [7]	✗	✗	12B	Collected	Python	Code: 159GB	✗	✗	HumanEval, APPS
AlphaCode [17]	✗	✓	41B	Collected	12 langs	Code: 715.1GB	✓	✗	HumanEval, APPS CodeContest
PaLM-Coder [8]	✗	✓	8B, 62B, 540B	Collected	Multiple	Text: 741B tokens Code: 39GB (780B tokens trained)	✓	✓	HumanEval, MBPP TransCoder, DeepFix
PolyCoder [41]	✓	✓	2.7B	Collected	12 langs	Code: 253.6GB	✗	✗	HumanEval
GPT-Neo [5]	✓	✓	1.3B, 2.7B	The Pile	Multiple	Text: 730GB Code: 96GB (400B tokens trained)	✗	✗	HumanEval
GPT-NeoX [4]	✓	✓	20B	The Pile	Multiple	Text: 730GB Code: 96GB (473B tokens trained)	✗	✗	HumanEval
GPT-J [38]	✓	✓	6B	The Pile	Multiple	Text: 730GB Code: 96GB (473B tokens trained)	✗	✗	HumanEval
InCoder [10]	✓	✓	1.3B, 6.7B	Collected	28 langs	Code: 159GB StackOverflow: 57GB (60B tokens trained)	✗	✗	HumanEval, MBPP CodeXGLUE
CodeGen-Multi [20]	✓	✓	6.1B, 16.1B	BigQuery	6 langs	Code: 150B tokens Text: 355B tokens (1000B tokens trained)	✗	✗	HumanEval, MTPB
CodeGen-Mono [20]	✓	✗	6.1B, 16.1B	BigPython	Python	Code: 150B tokens Text: 355B tokens (1300B tokens trained)	✗	✗	HumanEval, MTPB
CodeGeeX	✓	✓	13B	The Pile CodeParrot Collected	23 langs	Code: 158B tokens (850B tokens trained)	✓	✓	HumanEval-X, HumanEval MBPP, CodeXGLUE, XLCOST

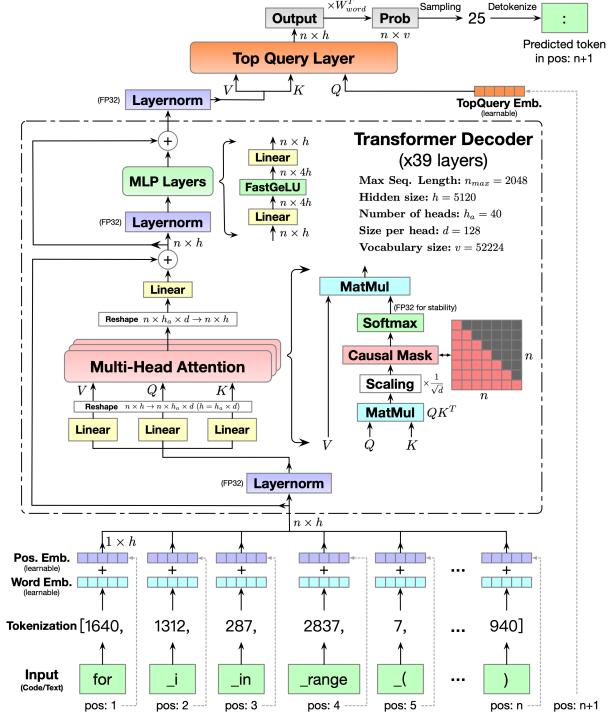


Figure 2: CodeGeeX’s model architecture. The model has 13B parameters, consisting of 39-layer left-to-right transformer decoders and a top query layer. It takes text/code tokens as input and outputs the probability of the next token autoregressively.

embedding of position $n + 1$. The final output is multiplied by the transpose of word embedding matrix to get the output probability. For decoding strategies, CodeGeeX supports greedy, temperature sampling, top-k sampling, top-p sampling, and beam search. Finally, detokenization will turn the selected token ID into an actual word.

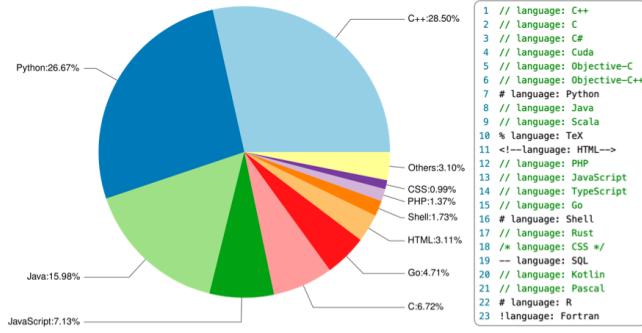


Figure 3: Language distribution and tags of CodeGeeX’s data.

2.2 Pre-Training Setup

Code Corpus. The training corpus contains two parts. The first part is from open source code datasets, the Pile [11] and CodeParrot¹. The Pile contains a subset of public repositories with more

¹<https://huggingface.co/datasets/transformersbook/codeparrot>

than 100 stars on GitHub, from which we select files of 23 popular programming languages including C++, Python, Java, JavaScript, C, Go, and so on. We identify the programming language of each file based on its suffix and the major language of the repository it belongs to. CodeParrot is another public Python dataset from BigQuery. The second part is supplementary data of Python, Java, and C++ directly scraped from GitHub public repositories that do not appear in the first part. We choose repositories that have at least one star and a total size within 10MB, then we filter out files that: 1) have more than 100 characters per line on average, 2) are automatically generated, 3) have a ratio of alphabet less than 40%, 4) are bigger than 100KB or smaller than 1KB. We format Python code according to the PEP8 standards.

Figure 3 shows the composition of the 158B-token training data, containing 23 programming languages. We divide the training data into segments of equal length. To help the model distinguish between multiple languages, we add a language-specific tag before each segment in the form of [Comment sign]language: [LANG], e.g., # language: Python.

Tokenization. The first step is to convert code snippets into numerical vectors. Considering that 1) there is a large number of natural language comments in code data, 2) the naming of variables, functions, and classes are often meaningful words, we treat code data in the same way as text data and apply the GPT-2 tokenizer [26]. It is a BPE (Byte Pair Encoding) [29] tokenizer that deals with the open-vocabulary problem using a fixed-size vocabulary with variable-length characters. The initial vocabulary size is 50,000, we encode multiple whitespaces as extra tokens following [7] to increase the encoding efficiency. Specifically, L whitespaces are represented by <|extratoken_X|>, where X=8+L. Since the vocabulary contains tokens from various natural languages, it allows CodeGeeX to process tokens in languages other than English, like Chinese, French, Russian, Japanese, and more. The final vocabulary size is $v = 52,224$. After tokenization, any code snippet or text description can be transformed into a vector of integers.

Word and Positional Embeddings. The next step is to associate each token with a word embedding. By looking up the token ID in a word embedding matrix $W_{word} \in \mathbb{R}^{v \times h}$, where $v = 52224$ is the vocabulary size (with extra tokens) and $h = 5120$ is the hidden size, a learnable embedding $x_{word} \in \mathbb{R}^h$ is obtained for each token. To capture positional information, we also adopt learnable positional embedding that maps the current position ID to a learnable embedding $x_{pos} \in \mathbb{R}^h$, from $W_{pos} \in \mathbb{R}^{n_{max} \times h}$, where $n_{max} = 2048$ is the maximum sequence length. Then, two embeddings are added to obtain the input embeddings $x_{in} = x_{word} + x_{pos}$ for the model. Finally, the entire sequence can be turned into input embeddings $X_{in} \in \mathbb{R}^{n \times h}$, where n is the input sequence length.

2.3 CodeGeeX Training

Parallel Training on Ascend 910. CodeGeeX was trained on a cluster of the Ascend 910 AI processors (32GB) with Mindspore (v1.7.0). We faced and addressed numerous unknown technical and engineering challenges during pre-training, as Ascend and Mindspore are relatively new compared to NVIDIA GPUs and PyTorch/TensorFlow. The entire pre-training process takes two

months on 192 nodes with 1,536 AI processors, during which the model consumes 850B tokens, equivalent to 5+ epochs (213,000 steps). Detailed configurations can be found in Table 2.

Table 2: Training configuration of CodeGeeX.

Category	Parameter	Value
Environment	Framework	Mindspore v1.7.0
	Hardwares	1,536x Ascend 910 AI processors
	Mem per GPU	32GB
	GPUs per node	8
	CPUs per node	192
Model	RAM per node	2048GB
	Model parameters	13B
	Vocabulary size	52224
	Position embedding	Learnable
	Maximum sequence length	2048
	Hidden size h	5120
	Feed-forward size $4h$	20480
	Feed-forward activation	FastGELU
	Layernorm epsilon	1e-5
	Layernorm precision	FP32
Parallelism	Number of attention heads h_n	40
	Attention softmax precision	FP32
	Dropout rate	0.1
	Model parallel size	8
	Data parallel size	192
Optimization	Global batch size	3072
	Optimizer	Adam
	Optimizer parameters	$\beta_1 = 0.9, \beta_2 = 0.999$
	Initial/final learning rate	1e-4/1e-6
	Warm-up step	2000
	Decay step	200000
	Learning rate scheduler	cosine decay
	Loss function \mathcal{L}	Cross entropy
	Loss scaling	Dynamic
	Loss scaling window	1000
	Trained steps	213000

To increase training efficiency, we adopt an 8-way model parallel training together with 192-way data parallel training, with ZeRO-2 [27] optimizer enabled to further reduce the memory consumption of optimizer states. Finally, the micro-batch size is 16 per node and the global batch size reaches 3,072.

Specifically, we use Adam optimizer [15] to optimize the loss in Equation 2. The model weights are under FP16, except that we use FP32 for layer-norm and softmax for higher precision and stability. The model takes ~27GB of GPU memory. We start from an initial learning rate 1e-4, and apply a cosine learning rate decay by:

$$lr_{current} = lr_{min} + 0.5 * (lr_{max} - lr_{min}) * (1 + \cos(\frac{n_{current}}{n_{decay}}\pi)) \quad (3)$$

Training Efficiency Optimization. Over the course of the training, we actively attempted to optimize the Mindspore framework to release the power of Ascend 910. Notably, we adopt the following techniques that significantly improve training efficiency:

- Kernel fusion: We fuse several element-wise operators to improve calculation efficiency on Ascend 910, including Bias+LayerNorm, BatchMatmul+Add, FastGELU+Matmul, Softmax, etc. We also optimize LayerNorm operator to support multi-core calculation.
- Auto Tune optimization²: When loading models, Mindspore first compiles them to static computational graphs. It uses the Auto

²<https://support.huawei.com/enterprise/en/doc/EDOC1100219270?section=j01g>

Tune tool to optimize the choice of operators (e.g., matrix multiplication in different dimensions). And it applies graph optimization techniques to deal with operator fusion and constant folding.

Table 3 shows the comparison of training efficiency before and after our optimization. The overall efficiency is measured by trained tokens per day. We observe that the efficiency per processor was improved 3× compared to the non-optimized implementation and the overall token throughput of 1,536 GPUs was improved by 224%.

Table 3: Training efficiency (before and after optimization).

	Before	After
Device	Ascend 910	Ascend 910
#GPUs	1536	1536
Parallelism	Data parallel + Model parallel	Data parallel + Model parallel
Sequence length	2048	2048
Global batch size	2048	3072
Step time(s)	15s	10s
Overall efficiency	24.2B tokens/day	54.3B tokens/day

2.4 Fast Inference

To serve the pre-trained CodeGeeX, we implement a pure PyTorch version of CodeGeeX that supports inference on NVIDIA GPUs. To achieve fast and memory-efficient inference, we apply both quantization and acceleration techniques to the pre-trained CodeGeeX.

Quantization. We apply post-training quantization techniques to decrease memory consumption of CodeGeeX during inference. We transform weights W in all linear transformations from FP16 to INT8 using the common absolute maximum quantization:

$$W_q = \text{Round}(\frac{W}{\lambda}), \lambda = \frac{\text{Max}(|W|)}{2^{b-1} - 1} \quad (4)$$

where b is the bitwidth and $b = 8$. λ is the scaling factor. This quantization transform FP16 values in $[-\text{Max}(|W|), \text{Max}(|W|)]$ to integers between $[-127, 127]$.

As in Table 4, the memory consumption of CodeGeeX decreases from ~26.9GB to ~14.7GB (down by 45.4%), allowing CodeGeeX inference on one RTX 3090 GPU. Importantly, Figure 4 shows that the quantization only slightly affects the performance on the code generation task (Cf. Section 3 for details about HumanEval-X.).

Table 4: GPU memory and inference time of CodeGeeX w/ and w/o quantization on different GPUs and frameworks.

Implementation	GPU	Format	L=128		L=512		L=2048	
			Mem (G)	Time (s)	Mem (G)	Time (s)	Mem (G)	Time (s)
Pytorch	A100	FP16	26.9	3.66	27.6	14.35	34.6	63.20
Pytorch	A100	INT8	14.7	9.40	16.1	37.38	18.7	155.01
Pytorch	3090	FP16			OOM			
Pytorch	3090	INT8	14.7	13.82	16.1	55.42	18.7	228.67
FastTrans	A100	FP16	26.0	2.43	26.3	10.21	27.5	50.09
FastTrans	A100	INT8	14.9	1.61	15.2	6.35	15.6	34.96
FastTrans	3090	FP16			OOM			
FastTrans	3090	INT8	14.5	2.25	14.8	9.34	16.0	43.81

Acceleration. After quantization, we further implement a faster version of CodeGeeX using FasterTransformer (FastTrans)³. It supports highly-optimized operations by using layer fusion, GEMM autotuning, and hardware-accelerated functions. For INT8 quantized version, we also implement a custom kernel that accelerates

³<https://github.com/NVIDIA/FasterTransformer>

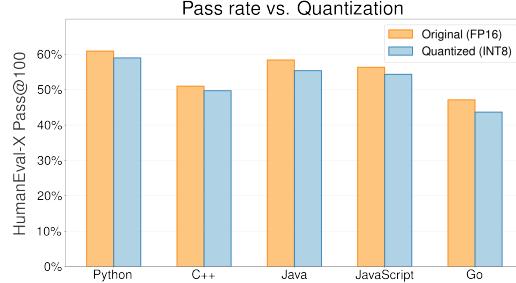


Figure 4: CodeGeeX vs. its quantized version on code generation of 5 programming languages in HumanEval-X.

the mixed precision matrix multiplication between INT8 weights and FP16 activation vectors. As in Table 4, the INT8 quantization plus FastTrans implementation achieves the fastest inference speed and the lowest GPU memory consumption on a single GPU. The inference time per token is within 13ms (1.61 seconds / 128 tokens).

3 THE HumanEval-X BENCHMARK

We develop the HumanEval-X benchmark⁴ for evaluating multilingual code models. There are 164 code problems defined for five major languages: C++, Java, JavaScript, Go, and Python, resulting in $164 \times 5 = 820$ problem-solution pairs. For each problem, it supports both code generation and code translation.

3.1 HumanEval-X: A Multilingual Benchmark

HumanEval [7] has been developed to evaluate Codex by OpenAI. However, similar to MBPP [2] and APPS [13], it only consists of handcrafted programming problems in Python, thus cannot be directly applied to systematically evaluate the performance of multilingual code generation. To this end, we propose to develop a multilingual variant of HumanEval, referred to as HumanEval-X. This is not trivial. For each problem, defined only for Python, in HumanEval, we manually rewrite its prompt, canonical solution, and test cases in the other four languages—C++, Java, JavaScript, and Go. Altogether, we have 820 problem-solution pairs in total in HumanEval-X, each comprising the following parts:

- **task_id:** programming language and numerical problem id, e.g., Java/0 represents the 0-th problem in Java;
- **declaration:** function declaration including necessary libraries or packages;
- **docstring:** description that specifies the functionality and example input/output;
- **prompt:** function declaration plus docstring;
- **canonical_solution:** a verified solution to the problem;
- **test:** test program including test cases.

Each problem-solution pair in HumanEval-X supports both code generation and code translation. An illustrative example is shown in Figure 5. We take the following efforts to make sure that the rewritten code conforms to the programming style of the corresponding language. First, we use the customary naming styles, like CamelCase in Java, Go, and JavaScript, and snake_case in C++. Second, we put the docstrings before the function declaration in Java, JavaScript,

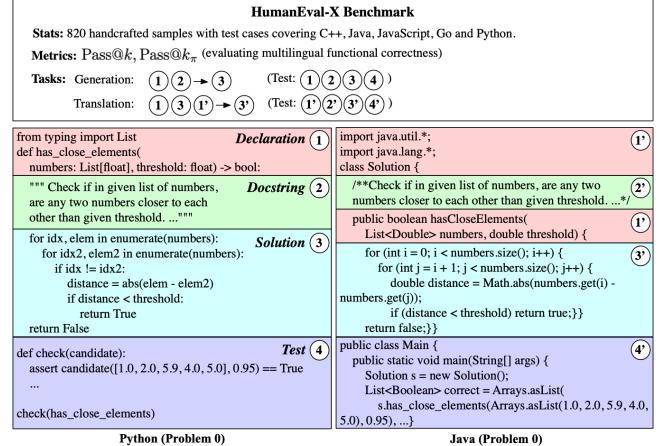


Figure 5: An illustration of HumanEval-X benchmark. Declarations, docstrings, solutions, and test cases are marked with red, green, blue, and purple respectively. *Generation* uses declaration and docstring as input to generate the solution. *Translation* uses declaration in both languages and solution in the source language as input, to generate solutions in the target language (docstring is not used to prevent models from directly solving the problem).

C++, and Go. Symbols in docstrings are modified, e.g., single quotes are replaced by double quotes in some languages, and keywords like True/False, None are also replaced. Third, we refine test cases according to language-specific behaviors, rather than forcing the programs to return the same result for different languages. For example, when converting an integer to a binary string, Python method bin adds a prefix “0b” before the string while Java method Integer.toBinaryString does not, so we remove such prefix in Java test cases. Last, we also take care of the rounding function. In Python, round converts half to the closest even number, unlike in other languages. Thus, we change the test cases to match the rounding implementations in each language.

3.2 HumanEval-X: Tasks

In HumanEval-X, we evaluate two tasks:

Code Generation. The task of code generation takes a problem description (e.g., “write a factorial function”) as input and generates the solution in the selected languages (Cf. Figure 1 (a)). Specifically, the model takes in the prompt including declaration and docstrings, and generates the implementation of the function. Note that HumanEval-X uses the same problem set for all the five languages, thus, for solving each problem, it supports either one single language or multiple languages simultaneously.

Code Translation. The task of code translation takes the implementation of a problem in the source language and generates its counterpart implementation in the target language. Precisely, its input includes the function declaration and a canonical solution in the source language (e.g., Python). The model should translate the solution to the target language. Adding declaration in the target language restricts function names and variable types, making the evaluation easier, especially under the zero-shot setting. To

⁴<https://hub.docker.com/r/codegeex/codegeex>

prevent the models from directly solving the problem rather than translating, we do not include the docstrings.

HumanEval-X supports the translation between all pairs of 5 languages, that is, in total 20 source-target language pairs.

Metric. For both tasks, we use test cases to evaluate the exact functional correctness of the generated code, measuring the performance with $\text{pass}@k$ [16]. Specifically, we use the unbiased method to estimate $\text{pass}@k$ [7]:

$$\text{pass}@k := \mathbb{E}[1 - \frac{\binom{n-c}{k}}{\binom{n}{k}}], n = 200, k \in \{1, 10, 100\} \quad (5)$$

where n is the total number of generations ($n=200$ in this work), k is the sampling budget (typically $k \in \{1, 10, 100\}$), and c is the number of samples that pass all test cases. We average over the problem set to get the expectation. $1 - \frac{\binom{n-c}{k}}{\binom{n}{k}}$ is the estimated $\text{pass}@k$ for a single problem, and \mathbb{E} is the expectation of $\text{pass}@k$ over all problems. In practice, we average single-problem $\text{pass}@k$ among all test-set problems to get the expectation.

Multilingual Metric with Budget Allocation. Unlike monolingual models, multilingual code models can solve problems by allocating generation budgets to various languages to increase the sampling diversity and improve the solve rate. Given a budget k , we can distribute part of it n_i to each language with the assignment

$$\pi = (n_1, n_2, \dots, n_m), \sum_{i=1}^m n_i = k, \quad (6)$$

where n_i is the generation budget assigned to language i , m is the number of candidate languages. Under an assignment $\pi = (n_1, \dots, n_m)$, for a problem p , the $\text{pass}@k_\pi$ can be estimated by:

$$\text{pass}@k_\pi = \mathbb{E}\left[1 - \prod_{i=1}^m \frac{\binom{n-c_i}{n_i}}{\binom{n}{n_i}}\right], \quad (7)$$

where n is the total number of generations, n_i is the sampling budget and c_i is the number of samples that pass all test cases for language i . We show in Section 4.2 that multilingual models can benefit from budget allocation strategies and have a higher solve rate than using any single language.

4 EVALUATING CodeGeeX ON HumanEval-X

We evaluate CodeGeeX for the code generation and translation tasks on the multilingual benchmark HumanEval-X. For baselines, we compare CodeGeeX with five competitive open-source baselines: GPT-J-6B [38], GPT-NeoX-20B [4], InCoder-6.7B [10], and CodeGen-Multi-6B/16B [20]. These models are all trained on multilingual code data, but is previously only evaluated in HumanEval (Python). And they are closer to the scale of CodeGeeX or even larger, while smaller models in the literature are ignored. For all baselines, we use the versions available on HuggingFace [40]. For each model, all $\text{pass}@k, k \in \{1, 10, 100\}$ results are estimated with $n = 200$. We follow the HumanEval-X's settings in Section 3.

4.1 Experimental Results

Multilingual Code Generation. Table 6 reports the code generation results in terms of the $\text{pass}@k, k \in \{1, 10, 100\}$ for CodeGeeX and five baseline models on HumanEval-X. CodeGeeX significantly

outperforms models trained with mixed corpora (GPT-J-6B and GPT-NeoX-20B), even though GPT-NeoX-20B has much more parameters. For models trained on codes, CodeGeeX outperforms those with smaller scales (InCoder-6.7B, CodeGen-Multi-6B) by a large margin, and is competitive with the larger CodeGen-Multi-16B model. CodeGeeX achieves the best average performance among all models, even slightly better than the larger CodeGen-Multi-16B in all three metrics (0.37%~1.67% improvements). When considering individual languages, models have preferences highly related to the training set distribution. For example, the best language for CodeGeeX is Python while for CodeGen-Multi-16B is Java.

Cross-Lingual Code Translation. Table 5 illustrates the results on code translation. For CodeGeeX, we evaluate both the original version CodeGeeX-13B and the fine-tuned CodeGeeX-13B-FT. CodeGeeX-13B-FT is first fine-tuned using the training set of code translation task in XLCoST [45], and then continuously fine-tuned by a small amount of Go data (since Go is missing in XLCoST). Among all translation pairs, CodeGeeX-13B-FT performs the best on $\text{pass}@100$ in 11 out of the 20, while CodeGen-Multi-16B is the best on 7 of them. We also observe a clear preference for languages by different models. CodeGeeX performs the best when translating other languages to Python and C++, while CodeGen-Multi-16B performs better when translating to JavaScript and Go.

4.2 Multilingual Understanding

We perform studies to understand whether and how multilingual pre-training can benefit problem-solving of CodeGeeX.

Exploration vs. Exploitation under Fixed Budgets. Given a fixed budget k , $\text{pass}@k$ evaluates the ability of models generating at least 1 correct solution under k generations. Previous works [7, 17] have already discovered that there's a trade-off between exploration and exploitation: When the budget is small, it is better to use a low temperature to ensure accuracy on easy problems. When the budget is large, instead, adjusting a higher temperature makes the model more likely to find at least one solution for difficult problems.

Pass Rate Distribution vs. Languages. Unlike monolingual models, multilingual models can solve problems more effectively using various programming languages. In Figure 6, we observe that the pass rate distribution of problems against different languages are diverse. This inspires us to use budget allocation methods to help improve the diversity of the generated solutions.

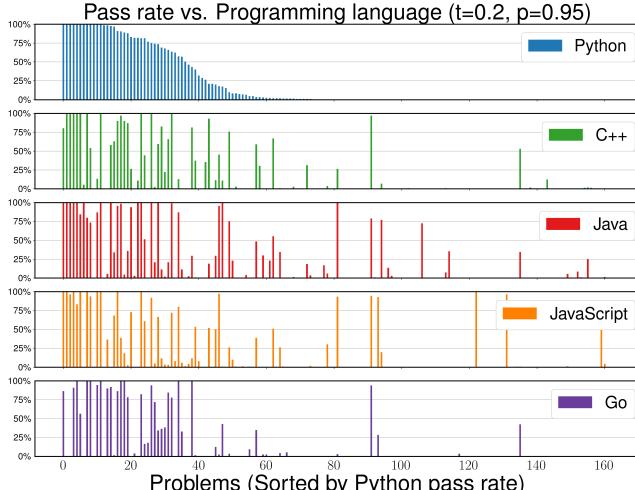
Budget Allocation Strategies. We compare three basic strategies: Best Single, choose a single language with the best performance; Uniform, allocate the budget uniformly; Weighted, allocate the budget to multiple languages based on their proportions in the training corpus (detailed weights can be found in Appendix Table 9). Table 7 illustrates how budget allocation improves multilingual generation. Both Uniform and Weighted outperform Best Single by promoting a more diverse generation, which gives a higher chance of solving problems. Weighted is slightly better due to the prior knowledge of the model. For model-wise comparison, CodeGeeX shows up a decent advantage over other baselines in both strategies, which suggests that it might have a more diverse solution set under multiple languages. In real-world scenarios, programming

Table 5: Code translation on HumanEval-X.

	Model	Target Language														
		Python				C++		Java				JavaScript			Go	
Py	InCoder-6.7B	-	-	-	26.11	41.00	54.25	26.74	42.66	61.20	37.05	58.85	78.91	15.69	27.57	43.67
	CodeGen-Multi-16B	-	-	-	35.94	47.81	59.37	29.27	45.70	64.45	43.40	66.26	82.55	28.87	41.01	57.72
	CodeGeeX-13B	-	-	-	26.54	43.56	56.48	25.84	41.52	59.72	23.22	47.33	65.87	9.56	23.83	33.56
	CodeGeeX-13B-FT	-	-	-	34.16	46.86	61.22	41.98	58.17	72.78	34.81	53.05	66.08	16.41	30.76	46.37
C++	InCoder-6.7B	34.37	58.41	78.57	-	-	-	34.04	57.02	68.70	37.05	65.05	79.61	25.54	39.11	58.02
	CodeGen-Multi-16B	33.83	55.37	76.64	-	-	-	43.20	69.84	88.82	54.51	71.50	83.14	27.94	49.73	68.32
	CodeGeeX-13B	27.18	49.02	67.69	-	-	-	22.56	40.91	64.08	30.23	55.68	75.58	8.64	18.79	31.76
	CodeGeeX-13B-FT	62.79	80.39	87.10	-	-	-	71.68	81.62	85.84	50.83	64.55	74.57	16.71	34.18	52.98
Java	InCoder-6.7B	42.76	65.55	80.43	40.01	55.17	70.39	-	-	-	43.20	68.24	84.39	21.58	35.20	54.97
	CodeGen-Multi-16B	52.73	69.30	82.74	41.42	54.68	65.50	-	-	-	57.65	67.90	79.22	34.00	48.49	67.94
	CodeGeeX-13B	43.41	68.46	84.03	39.33	58.48	72.36	-	-	-	44.19	64.22	82.89	17.17	32.74	47.71
	CodeGeeX-13B-FT	75.03	87.71	95.13	49.67	65.65	75.40	-	-	-	49.95	62.82	79.64	18.85	32.92	48.93
JS	InCoder-6.7B	23.18	50.47	67.26	35.47	54.48	70.71	30.67	50.90	71.03	-	-	-	25.79	42.96	61.47
	CodeGen-Multi-16B	35.52	52.23	69.78	35.41	53.12	64.47	33.79	56.06	74.00	-	-	-	33.38	49.08	64.14
	CodeGeeX-13B	31.15	54.02	72.36	30.32	51.63	69.37	24.68	48.35	69.03	-	-	-	11.91	26.39	39.81
	CodeGeeX-13B-FT	67.63	81.88	89.30	46.87	60.82	73.18	56.55	70.27	80.71	-	-	-	16.46	32.99	50.29
Go	InCoder-6.7B	34.14	54.52	70.88	30.45	48.47	62.81	34.52	53.95	69.92	39.37	63.63	80.75	-	-	-
	CodeGen-Multi-16B	38.32	50.57	68.65	32.95	45.88	59.56	36.55	59.12	78.70	38.93	56.68	70.68	-	-	-
	CodeGeeX-13B	35.92	56.02	77.32	29.83	41.98	58.15	22.89	41.04	61.46	25.24	46.50	69.93	-	-	-
	CodeGeeX-13B-FT	57.98	79.04	93.57	38.97	53.05	63.92	54.22	69.03	79.40	43.07	59.78	74.04	-	-	-

Table 6: Code generation on HumanEval-X.

Language	Metric	GPT-J -6B	GPT-NeoX -20B	InCoder -6.7B	CodeGen -Multi-6B	CodeGen -Multi-16B	CodeGeeX -13B (ours)
Python (HumanEval)	pass@1	11.10%	13.83%	16.41%	19.41%	19.22%	22.89%
	pass@10	18.67%	22.72%	26.55%	30.29%	34.64%	39.57%
	pass@100	30.98%	39.56%	43.95%	49.63%	55.17%	60.92%
C++	pass@1	7.54%	9.90%	9.50%	11.44%	18.05%	17.06%
	pass@10	13.67%	18.99%	19.30%	26.23%	30.84%	32.21%
	pass@100	30.16%	38.75%	36.10%	42.82%	50.90%	51.00%
Java	pass@1	7.86%	8.87%	9.05%	15.17%	14.95%	20.04%
	pass@10	14.37%	19.55%	18.64%	31.74%	36.73%	36.70%
	pass@100	32.96%	42.23%	40.70%	53.91%	60.62%	58.42%
JavaScript	pass@1	8.99%	11.28%	12.98%	15.41%	18.40%	17.59%
	pass@10	16.32%	20.78%	22.98%	27.92%	32.80%	32.28%
	pass@100	33.77%	42.67%	43.34%	48.81%	56.48%	56.33%
Go	pass@1	4.01%	5.00%	8.68%	9.98%	13.03%	14.43%
	pass@10	10.81%	15.70%	13.80%	23.26%	25.46%	25.68%
	pass@100	23.70%	32.08%	28.31%	41.01%	48.77%	47.14%
Average	pass@1	7.90%	9.78%	11.33%	14.28%	16.73%	18.40%
	pass@10	14.77%	19.55%	20.25%	27.89%	32.09%	33.29%
	pass@100	30.32%	39.06%	38.48%	47.24%	54.39%	54.76%

**Figure 6: In HumanEval-X, each problem's pass rate varies when generating in different programming languages with CodeGeeX. Left: $t = 0.2, p = 0.95$; Right: $t = 0.8, p = 0.95$.****Table 7: Results for fixed-budget multilingual generation on HumanEval-X.** Best model-wise performances on methods are **bolded**, while best method-wise performances are in *italic*.

Metric	Method	GPT-J -6B	GPT-NeoX -20B	InCoder -6.7B	CodeGen -Multi-6B	CodeGen -Multi-16B	CodeGeeX -13B
pass@ k_{π} ($k = 100$)	Best Single	33.77%	42.67%	43.95%	53.19%	60.62%	60.92%
	Uniform	36.40%	44.75%	43.89%	53.47%	61.01%	62.41%
	Weighted	36.76%	44.97%	45.60%	53.94%	61.34%	62.95%

languages are created with a specific purpose and unique design. With the proper budget allocation strategy, we can take advantage of the model's multilingual ability for specific tasks.

Test Result Characteristics. To study how models actually behave on programming problems, we group the generated samples' test results into five categories: Passed, Wrong Answer, Runtime Error, Syntax/Semantic Error and Unfinished. More precisely, Runtime Error includes an out-of-bound index, wrong string format, etc; Syntax/Semantic Error indicates errors detected by syntax or semantic check, like compilation error in compiled languages and syntax/undefined/type error in interpreted languages; Unfinished means that the model fails to complete one function within maximum length. As in Figure 7, the most common error type is Wrong Answer, with a ratio ranging from 0.44 to 0.75 (except for Go), showing that code generation models at the current stage mainly suffer from incorrect code logic rather than semantics. Models have a high syntax error rate with Go, which may be due to Go's strict restrictions on syntax, *i.e.*, forbidding unused variables and imports, thus failing to compile many logically correct codes. Overall, CodeGeeX is less probable to generate code that has Runtime or Syntax/Semantic Error.

Negative Correlation in Translation. When evaluating the translation ability in HumanEval-X, an interesting observation is that the performance of A-to-B and B-to-A are usually negatively correlated, shown in Figure 8. Such asymmetry suggests that multilingual code generation models may have an imbalanced focus on source and target languages during code translation. We provide two possible explanations. First, language distributions in the training corpus

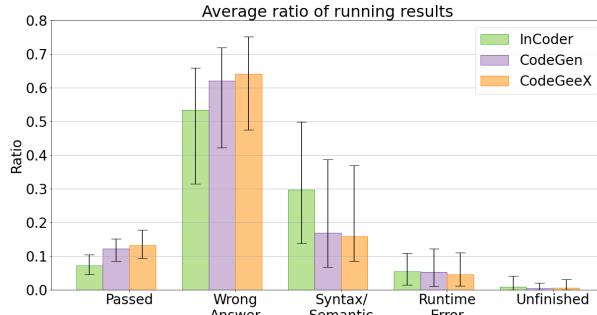


Figure 7: Test result statistics across models. For each model and language, we study 200 samples generated under $t = 0.8$, $p = 0.95$. CodeGeeX has less Runtime or Syntax/Semantic Error.

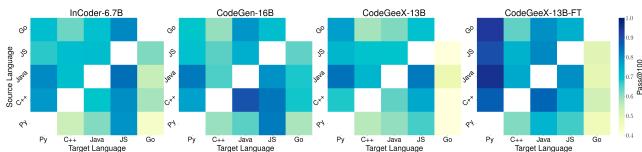


Figure 8: Performance of translating A-to-B is negatively correlated with B-to-A. Such asymmetry indicates that multilingual models lack a high-level understanding between languages.

differ a lot, resulting in different levels of generation ability. For example, the ratio of Python is 26.6% (vs. Go 4.7%) in CodeGeeX training corpus, and average pass@100 of *Others-to-Python* reaches ~90% (vs. *Others-to-Go* only ~50%). Second, some languages are themselves harder to automatically write with syntactic and semantic accuracy due to language-dependent features, affecting translation performance as target languages. For instance, Go, which models translate poorly into, has more constraints on the syntax level, like forbidding unused variables or imports.

5 THE CODEGEEEX TOOLS AND USERS

Based on CodeGeeX, we build open-source extensions for IDEs including VS Code, JetBrains and Cloud Studio. The extensions support code generation, completion, translation and explanation, aiming at improving the development efficiency of programmers. As of this writing, CodeGeeX has served tens of thousands of users, with an average of 200+ API calls per active user per weekday.

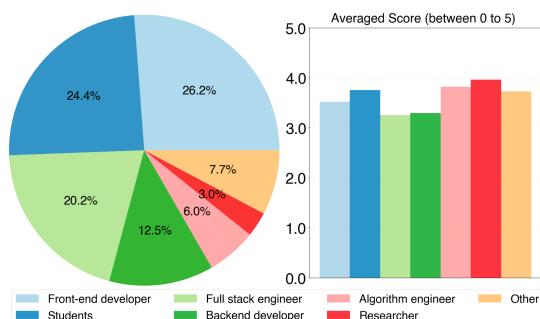


Figure 9: Profession vs. satisfaction. Left: Profession distribution. Right: Averaged rating score of CodeGeeX extensions.

We took a survey on CodeGeeX's user experience from 168 users covering *front-end developer*, *backend developer*, *full stack engineer*, *algorithm engineer*, *students*, *researcher*, and *other programmers*. Figure 9 illustrates users' profession distribution and the satisfaction score. We evaluate the satisfaction considering five dimensions, "Ease of Use", "Reliability", "Feature", "Visual", "Speed", each scored from 0 to 5. Figure 9 shows that the majority of users have positive experiences with CodeGeeX, especially for researchers and students.

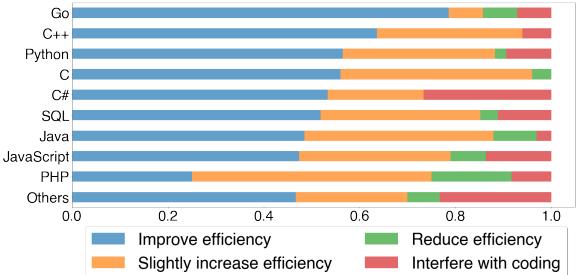


Figure 10: Survey on "Has CodeGeeX improved your coding efficiency?" Over 83.4% of users have positive answers.

We further investigate how multilinguality of CodeGeeX help coding. Figure 10 illustrates how users evaluate the helpfulness of CodeGeeX during development. There are on average over 83.4% of users think CodeGeeX can improve or slightly increase their coding efficiency, especially for mainstream programming languages like Go, C++, Python, C, C#, etc. Note that these well-performing programming languages also appear more frequently in the training data (Figure 3), which encourages us to train CodeGeeX on more language-specific data to enhance its capability.

6 CONCLUSION

We introduce CodeGeeX, a 13B pre-trained 23-language code generation model, as well as we build HumanEval-X, to fill the gap of multilingual code generation. CodeGeeX consistently outperforms open-sourced multilingual baselines of the same scale on code generation and translation tasks. The extensions built on CodeGeeX bring significant benefits in increasing coding efficiency. The multilinguality of CodeGeeX brings the potential of solving problems with a ubiquitous set of formalized languages. We open sourced CodeGeeX aiming to help researchers and developers to widely take benefit of large pre-trained models for code generation.

7 ACKNOWLEDGMENTS

This research was supported by Natural Science Foundation of China (NSFC) for Distinguished Young Scholars No. 61825602, NSFC No. 62276148 and No. 61836013, and a research fund from Zhipu.AI. We give our special thanks to Wenguang Chen from Tsinghua, the Peng Cheng Laboratory, and Zhipu.AI for sponsoring the training and inference GPU resources. We thank all our collaborators and partners from Tsinghua KEG, IIIS, Peng Cheng Laboratory, and Zhipu.AI, including Aohan Zeng, Wendi Zheng, Lilong Xue, Yifeng Liu, Yanru Chen, Yichen Xu, Qingyu Chen, Zhongqi Li, Gaojun Fan, Yifan Yao, Qihui Deng, Bin Zhou, Ruijie Cheng, Peinan Yu, Jingyao Zhang, Bowen Huang, Zhaoyu Wang, Jiecai Shan, Xuyang Ding, Xuan Xue, and Peng Zhang.

REFERENCES

- [1] Wasi Uddin Ahmad, Saikat Chakraborty, Baishakhi Ray, and Kai-Wei Chang. 2021. Unified pre-training for program understanding and generation. *arXiv preprint arXiv:2103.06333* (2021).
- [2] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. 2021. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732* (2021).
- [3] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. 2016. Layer normalization. *arXiv preprint arXiv:1607.06450* (2016).
- [4] Sid Black, Stella Biderman, Eric Hallahan, Quentin Anthony, Leo Gao, Laurence Golding, Horace He, Connor Leahy, Kyle McDonell, Jason Phang, et al. 2022. Gpt-neox-20b: An open-source autoregressive language model. *arXiv preprint arXiv:2204.06745* (2022).
- [5] Sid Black, Leo Gao, Phil Wang, Connor Leahy, and Stella Biderman. 2021. GPT-Neo: Large Scale Autoregressive Language Modeling with Mesh-Tensorflow. *If you use this software, please cite it using these metadata* 58 (2021).
- [6] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems* 33 (2020), 1877–1901.
- [7] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374* (2021).
- [8] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, and et al. 2022. Palm: Scaling language modeling with pathways. <https://arxiv.org/abs/2204.02311>
- [9] Zhangyin Feng, Daya Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Dixin Jiang, et al. 2020. Codebert: A pre-trained model for programming and natural languages. *arXiv preprint arXiv:2002.08155* (2020).
- [10] Daniel Fried, Armen Aghajanyan, Jessy Lin, Sida Wang, Eric Wallace, Freda Shi, Ruiqi Zhong, Wen-tau Yih, Luke Zettlemoyer, and Mike Lewis. 2022. Incoder: A generative model for code infilling and synthesis. *arXiv preprint arXiv:2204.05999* (2022).
- [11] Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, et al. 2020. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027* (2020).
- [12] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 770–778.
- [13] Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, et al. 2021. Measuring coding challenge competence with apps. *arXiv preprint arXiv:2105.09938* (2021).
- [14] Dan Hendrycks and Kevin Gimpel. 2016. Gaussian error linear units (gelus). *arXiv preprint arXiv:1606.08415* (2016).
- [15] Diederik P Kingma and Jimmy Ba. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980* (2014).
- [16] Sumith Kulal, Panupong Pasupat, Kartik Chandra, Mina Lee, Oded Padon, Alex Aiken, and Percy S Liang. 2019. SPoC: Search-based Pseudocode to Code. In *Advances in Neural Information Processing Systems*, H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and R. Garnett (Eds.), Vol. 32. Curran Associates, Inc. <https://proceedings.neurips.cc/paper/2019/file/7298332f04ac004a0ca44cc69ecfcf6b-Paper.pdf>
- [17] Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittweis, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, et al. 2022. Competition-Level Code Generation with AlphaCode. *arXiv preprint arXiv:2203.07814* (2022).
- [18] Shuai Lu, Daya Guo, Shuo Ren, Junjie Huang, Alexey Svyatkovskiy, Ambrosio Blanco, Colin Clement, Dawn Drain, Dixin Jiang, Duyu Tang, et al. 2021. Codexglue: A machine learning benchmark dataset for code understanding and generation. *arXiv preprint arXiv:2102.04664* (2021).
- [19] Lili Mou, Rui Men, Ge Li, Lu Zhang, and Zhi Jin. 2015. On end-to-end program generation from user intention by deep neural networks. *arXiv preprint arXiv:1510.07211* (2015).
- [20] Erik Nijkamp, Bo Pang, Hiroaki Hayashi, Lifu Tu, Huan Wang, Yingbo Zhou, Silvio Savarese, and Caiming Xiong. 2022. A conversational paradigm for program synthesis. *arXiv preprint arXiv:2203.13474* (2022).
- [21] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*. 311–318.
- [22] Long Phan, Hieu Tran, Daniel Le, Hieu Nguyen, James Anibal, Alec Peltekian, and Yanfang Ye. 2021. Cotext: Multi-task learning with code-text transformer. *arXiv preprint arXiv:2105.08645* (2021).
- [23] Oleksandr Polozov and Sumit Gulwani. 2015. Flashmeta: A framework for inductive program synthesis. In *Proceedings of the 2015 ACM SIGPLAN International Conference on Object-Oriented Programming, Systems, Languages, and Applications*. 107–126.
- [24] Weizhen Qi, Yeyun Gong, Yu Yan, Can Xu, Bolun Yao, Bartuer Zhou, Biao Cheng, Daxin Jiang, Jiusheng Chen, Ruofei Zhang, et al. 2021. ProphetNet-x: large-scale pre-training models for English, Chinese, multi-lingual, dialog, and code generation. *arXiv preprint arXiv:2104.08006* (2021).
- [25] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre training. (2018).
- [26] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. 2019. Language models are unsupervised multitask learners. *OpenAI blog* 1, 8 (2019), 9.
- [27] Samyam Rajbhandari, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. 2020. Zero: Memory optimizations toward training trillion parameter models. In *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*. IEEE, 1–16.
- [28] Shuo Ren, Daya Guo, Shuai Lu, Long Zhou, Shujie Liu, Duyu Tang, Neel Sundaresan, Ming Zhou, Ambrosio Blanco, and Shuai Ma. 2020. Codebleu: a method for automatic evaluation of code synthesis. *arXiv preprint arXiv:2009.10297* (2020).
- [29] Rico Sennrich, Barry Haddow, and Alexandra Birch. 2015. Neural machine translation of rare words with subword units. *arXiv preprint arXiv:1508.07090* (2015).
- [30] Armando Solar-Lezama. 2008. *Program synthesis by sketching*. University of California, Berkeley.
- [31] Phillip D Summers. 1977. A methodology for LISP program construction from examples. *Journal of the ACM (JACM)* 24, 1 (1977), 161–175.
- [32] Zeyu Sun, Qihao Zhu, Yingfei Xiong, Yican Sun, Lili Mou, and Lu Zhang. 2020. Treegen: A tree-based transformer architecture for code generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 34. 8984–8991.
- [33] Alexey Svyatkovskiy, Shaokun Deng, Shengyu Fu, and Neel Sundaresan. 2020. Intellicode compose: Code generation using transformer. In *Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. 1433–1443.
- [34] Romal Thoppilan, Daniel De Freitas, Jamie Hall, Noam Shazeer, Apoorv Kulshreshtha, Heng-Tze Cheng, Alicia Jin, Taylor Bos, Leslie Baker, Yu Du, et al. 2022. Lamda: Language models for dialog applications. *arXiv preprint arXiv:2201.08239* (2022).
- [35] Lewis Tunstall, Leandro von Werra, and Thomas Wolf. 2022. *Natural language processing with transformers*. "O'Reilly Media, Inc."
- [36] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is All you Need. In *Advances in Neural Information Processing Systems*, I. Guyon, U. Von Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett (Eds.), Vol. 30. Curran Associates, Inc. <https://proceedings.neurips.cc/paper/2017/file/3f5ee243547dee91fbd053c1e4a845aa-Paper.pdf>
- [37] Richard J Walddinger and Richard CT Lee. 1969. PROW: A step toward automatic program writing. In *Proceedings of the 1st international joint conference on Artificial intelligence*. 241–252.
- [38] Ben Wang and Aran Komatsuzaki. 2021. GPT-J-6B: A 6 Billion Parameter Autoregressive Language Model. <https://github.com/kingoflolz/mesh-transformer-jax>
- [39] Yu Wang, Weishi Wang, Shafiq Joty, and Steven CH Hoi. 2021. Codet5: Identifier-aware unified pre-trained encoder-decoder models for code understanding and generation. *arXiv preprint arXiv:2109.00859* (2021).
- [40] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierrick Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. 2019. Huggingface’s transformers: State-of-the-art natural language processing. *arXiv preprint arXiv:1910.03771* (2019).
- [41] Frank F Xu, Uri Alon, Graham Neubig, and Vincent Josua Hellendoorn. 2022. A systematic evaluation of large language models of code. In *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming*. 1–10.
- [42] Aohan Zeng, Xiao Liu, Zhengxiao Du, Zihan Wang, Hanyu Lai, Ming Ding, Zhuoyi Yang, Yifan Xu, Wendi Zheng, Xiao Xia, et al. 2022. Clm-130b: An open bilingual pre-trained model. *arXiv preprint arXiv:2210.02414* (2022).
- [43] Wei Zeng, Xiaozhe Ren, Teng Su, Hui Wang, Yi Liao, Zhiwei Wang, Xin Jiang, ZhenZhang Yang, Kaisheng Wang, Xiaoda Zhang, Chen Li, Ziyan Gong, Yifan Yao, Xinjing Huang, Jun Wang, Jianfeng Yu, Qi Guo, Yue Yu, Yan Zhang, Jin Wang, Hengtao Tao, DASEN Yan, ZEXUAN Yi, Fang Peng, Fangqing Jiang, Han Zhang, Lingfeng Deng, Yehong Zhang, Zhe Lin, Chao Zhang, Shaojie Zhang, Mingyue Guo, Shanzhi Gu, Gaojun Fan, Yaowei Wang, Xuefeng Jin, Qun Liu, and Yonghong Tian. 2021. PanGu- α : Large-scale Autoregressive Pretrained Chinese Language Models with Auto-parallel Computation. (2021).
- [44] Susan Zhang, Stephen Roller, Namati Goyal, Mikel Artetxe, Moya Chen, Shuhui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. 2022. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068* (2022).

- [45] Ming Zhu, Aneesh Jain, Karthik Suresh, Roshan Ravindran, Sindhu Tipirneni, and Chandan K Reddy. 2022. XLCOST: A Benchmark Dataset for Cross-lingual Code Intelligence. *arXiv preprint arXiv:2206.08474* (2022).
- [46] Albert Ziegler, Eirini Kalliamvakou, X Alice Li, Andrew Rice, Devon Rifkin, Shawn Simister, Ganesh Sittampalam, and Edward Afandilian. 2022. Productivity assessment of neural code completion. In *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming*. 21–29.

A APPENDIX

A.1 Statistics of Code Corpus

Table 8 summarizes the composition of CodeGeeX’s code corpus.

Table 8: Composition of our code corpus for pre-training.

Language	# Tokens (B)	% Tokens (%)	Language Tag
C++	45.2283	28.4963	// language: C++
Python	42.3250	26.67	# language: Python
Java	25.3667	15.9824	// language: Java
JavaScript	11.3165	7.13	// language: JavaScript
C	10.6590	6.7157	// language: C
Go	7.4774	4.7112	// language: Go
HTML	4.9355	3.1096	<!--language: HTML-->
Shell	2.7498	1.7325	# language: Shell
PHP	2.1698	1.3671	// language: PHP
CSS	1.5674	0.9876	/* language: CSS */
TypeScript	1.1667	0.7351	// language: TypeScript
SQL	1.1533	0.7267	- language: SQL
TeX	0.8257	0.5202	% language: TeX
Rust	0.5228	0.3294	// language: Rust
Objective-C	0.4526	0.2851	// language: Objective-C
Scala	0.3786	0.2385	// language: Scala
Kotlin	0.1707	0.1075	// language: Kotlin
Pascal	0.0839	0.0529	// language: Pascal
Fortran	0.077	0.0485	!language: Fortran
R	0.0447	0.0281	# language: R
Cuda	0.0223	0.014	// language: Cuda
C#	0.0218	0.0138	// language: C#
Objective-C++	0.0014	0.0009	// language: Objective-C++

A.2 Details of Budget Allocation Strategies

As in Table 9, we allocate the budget to multiple languages based on their proportions in the training corpus.

Table 9: Detailed assignment of budget allocation strategies.

Strategy	Model	Python	C++	Java	JavaScript	Go
Uniform	All	20	20	20	20	20
Weighted	GPT-J-6B	17	36	11	22	14
	GPT-NeoX-20B	17	36	11	22	14
	InCoder-6.7B	45	12	5	34	4
	CodeGen-Multi-6B/16B	17	38	29	8	8
	CodeGeeX-13B (ours)	32	33	20	9	6

A.3 Evaluation on other benchmarks

A.3.1 Evaluation on HumanEval. The evaluation setting on HumanEval is the same as HumanEval-X. We show that among multilingual code generation models, CodeGeeX achieves the second highest performance on HumanEval, reaching 60% in pass@100 (surpassed by PaLMCoder-540B). We also notice that monolingual models outperform multilingual ones by a large margin, indicating that multilingual models might require a larger model capacity to master different languages.

Table 10: The results of CodeGeeX on HumanEval.

Model	Size	Type	Available	pass@1	pass@10	pass@100
CodeParrot [35]	1.5B	Multi	Yes	4.00%	8.70%	17.90%
PolyCoder [41]	2.7B	Multi	Yes	5.60%	9.80%	17.70%
GPT-J [38]	6B	Multi	Yes	11.60%	15.70%	27.70%
CodeGen-Multi [20]	6.1B	Multi	Yes	18.16%	27.81%	44.85%
InCoder [10]	6.7B	Multi	Yes	15.20%	27.80%	47.00%
GPT-NeoX [4]	20B	Multi	Yes	15.40%	25.60%	41.20%
LaMDA [34]	137B	Multi	No	14.00%*	-	47.30%*
CodeGen-Multi [20]	16.1B	Multi	Yes	19.22%	34.64%	55.17%
PaLM-Coder [8]	540B	Multi	No	36.00%*	-	88.40%*
Codex [7]	12B	Mono	No	28.81%	46.81%	72.31%
CodeGen-Mono [20]	16.1B	Mono	Yes	29.28%	49.86%	75.00%
CodeGeeX (ours)	13B	Multi	Yes	22.89%	39.57%	60.92%

A.3.2 Evaluation on MBPP. MBPP dataset is proposed by [2], containing 974 problems in Python. Due to specific input-output format, MBPP need to be evaluated under a few-shot setting. We follow the splitting in the original paper and use problems 11-510 for testing. Under 1-shot setting, we use problem 2 in prompts. Under 3-shot setting, we use problem 2,3,4 in prompts. The metric is pass@ k , $k \in \{1, 10, 80\}$. For pass@1, the temperature is 0.2 and top-p is 0.95; for pass@10 and pass@80, the temperature is 0.8 and top-p is 0.95. For baselines, we consider LaMDA-137B, PaLM-540B, Davinci-Codex (online API version of OpenAI Codex), PaLMCoder-540B and InCoder-6.7B.

The results indicate that the model capacity is essential for multilingual code generation model. With significantly more parameters, PaLM and Codex outperform CodeGeeX with a large margin. Meanwhile, we find that more shot in the prompts harm the performance of CodeGeeX, the same phenomenon have also been discovered in InCoder [10]. We assume that it is because smaller models do not have enough reasoning ability to benefit from the few-shot setting.

Table 11: The results of CodeGeeX on MBPP dataset [2].

Method	Model	Pass@1	Pass@10	Pass@80
3-shot	LaMDA-137B [2]	14.80	-	62.40
	PaLM-540B [8]	36.80	-	75.00
	Davinci-Codex [7]	50.40	-	84.40
	PaLMCoder-540B [8]	47.00	-	80.80
	CodeGeeX-13B (ours)	22.44	43.24	63.52
1-shot	InCoder-6.7B [10]	19.40	-	-
	CodeGeeX-13B (ours)	24.37	47.95	68.50

A.3.3 Evaluation on CodeXGLUE. CodeXGLUE [18] contains multiple datasets to support evaluation on multiple tasks, using similarity-based metrics like CodeBLEU, BLEU, and accuracy for generation tasks. We test the performance of CodeGeeX on the **code summarization** task. We first fine-tune CodeGeeX by mixing the training data in all languages to get one fine-tuned model. Then, we test the performance of the fine-tuned model in each language, using the BLEU score for evaluation because the models generate natural language in summarization tasks.

For all languages, we set the temperature to 0.2 and top-p to 0.95, and generate one summarization for each sample in the test set. We report the results in Table 14. CodeGeeX obtains an average BLEU score of 20.63, besting all baseline models. It is worth noting that after removing the results on Ruby (that CodeGeeX is

Table 12: The results of CodeGeeX on code translation in XLCoST benchmark. Six languages are considered, C++, Java, Python, C#, JavaScript, PHP, C. The metric is CodeBLEU [28]. The results of baselines are adopted from the original paper [45].

	Model	Snippet-level							Program-level						
		C++	Java	Py	C#	JS	PHP	C	C++	Java	Py	C#	JS	PHP	C
C++	CodeBERT	–	84.94	74.55	84.99	82.79	68.56	45.46	–	74.73	24.96	76.35	72.95	50.40	21.84
	PLBART	–	83.85	74.89	84.57	83.19	68.62	83.95	–	75.26	70.13	78.01	61.85	67.01	72.59
	CodeT5	–	86.35	76.28	85.85	84.31	69.87	90.45	–	80.03	71.56	81.73	79.48	70.44	85.67
	CodeGeeX	–	86.99	74.73	86.63	84.83	70.30	94.04	–	84.40	73.89	84.49	82.20	71.18	87.32
Java	CodeBERT	87.27	–	58.39	92.26	84.63	67.26	39.94	79.36	–	8.51	84.43	76.02	51.42	21.22
	PLBART	87.31	–	58.30	90.78	85.42	67.44	72.47	81.41	–	66.29	83.34	80.14	67.12	63.37
	CodeT5	88.26	–	74.59	92.56	86.22	69.02	82.78	84.26	–	69.57	87.79	80.67	69.44	78.78
	CodeGeeX	89.08	–	74.65	92.94	86.96	69.77	88.44	87.07	–	73.11	91.78	84.34	70.61	81.07
Py	CodeBERT	80.46	58.50	–	54.72	57.38	65.14	10.70	68.87	28.22	–	17.80	23.65	49.30	18.32
	PLBART	80.15	74.15	–	73.50	73.20	66.12	62.15	74.38	67.80	–	66.03	69.30	64.85	29.05
	CodeT5	81.56	78.61	–	78.89	77.76	67.54	68.67	78.85	73.15	–	73.35	71.80	67.50	56.35
	CodeGeeX	82.91	81.93	–	81.30	79.83	67.99	82.59	82.49	79.03	–	80.01	77.47	68.91	71.67
C#	CodeBERT	86.96	90.15	56.92	–	84.38	67.18	40.43	78.52	82.25	10.82	–	75.46	51.76	21.63
	PLBART	84.98	6.27	69.82	–	85.02	67.30	75.74	80.17	81.37	67.02	–	79.81	67.12	57.60
	CodeT5	88.06	91.69	73.85	–	85.95	68.97	81.09	83.59	85.70	69.52	–	80.50	69.63	77.35
	CodeGeeX	88.70	93.03	74.55	–	86.44	69.49	86.69	87.11	90.46	72.89	–	83.83	70.58	80.73
JS	CodeBERT	84.38	84.42	52.57	84.74	–	66.66	33.29	75.43	72.33	9.19	75.47	–	52.08	19.79
	PLBART	84.45	84.90	69.29	85.05	–	67.09	72.65	80.19	76.96	64.18	78.51	–	67.24	67.70
	CodeT5	85.06	85.48	73.15	85.96	–	68.42	80.49	82.14	79.91	68.42	81.77	–	68.76	74.57
	CodeGeeX	86.72	86.96	73.25	86.41	–	69.00	83.85	85.84	83.85	72.11	85.35	–	69.80	79.41
PHP	CodeBERT	82.58	81.57	69.29	80.96	79.94	–	28.45	50.13	46.81	16.92	49.75	48.12	–	22.19
	PLBART	83.87	81.66	71.17	78.00	82.94	–	57.39	79.40	72.77	61.26	74.16	44.26	–	56.23
	CodeT5	86.33	85.12	73.22	84.56	83.56	–	79.30	85.55	82.09	72.26	83.79	81.72	–	65.86
	CodeGeeX	86.75	86.24	71.37	85.58	84.17	–	83.89	87.23	83.90	71.02	85.34	82.81	–	78.76
C	CodeBERT	45.84	39.69	13.55	39.71	29.85	38.88	–	21.70	21.27	21.10	19.50	15.64	31.71	–
	PLBART	82.53	72.35	49.16	75.78	75.05	60.86	–	78.42	13.45	5.53	45.15	31.47	25.17	–
	CodeT5	90.26	81.81	63.81	83.05	79.73	66.32	–	88.17	76.12	56.32	80.20	76.50	64.28	–
	CodeGeeX	91.30	85.58	71.52	87.52	84.91	68.52	–	88.21	82.46	69.78	85.56	81.21	68.80	–

Table 13: The results of CodeGeeX on code summarization task in CodeXGLUE [18].

Model	All	Ruby	JavaScript	Go	Python	Java	PHP
CodeBERT [9]	17.83	12.16	14.90	18.07	19.06	17.65	25.16
PLBART [1]	18.32	14.11	15.56	18.91	19.30	18.45	23.58
ProphetNet-X [24]	18.54	14.37	16.60	18.43	17.87	19.39	24.57
CoTexT [22]	18.55	14.02	14.96	18.86	19.73	19.06	24.68
PolyglotCodeBERT [9]	19.06	14.75	15.80	18.77	18.71	20.11	26.23
DistillCodeT5 [39]	20.01	15.75	16.42	20.21	20.59	20.51	26.58
CodeGeeX (ours)	20.63	10.05*	16.01	24.62	22.50	19.60	31.00

not trained on), CodeGeeX outperforms the best baseline model (DistillCodeT5 [39]) by 1.88 in the average BLEU score.

A.3.4 Evaluation on XLCoST. XLCoST is a benchmark proposed by [45], containing parallel multilingual code data, with code snippets aligned among different languages. For generation tasks, XLCoST uses CodeBLEU, BLEU for evaluation. We choose the **code translation** task of XLCoST for CodeGeeX evaluation. We first fine-tune the parameters of CodeGeeX on the given training set, combining the training data in all 42 language pairs to obtain one fine-tuned

Table 14: The results of CodeGeeX on code summarization in CodeXGLUE benchmark [18]. Six languages are considered, Ruby, JavaScript, Go, Python, Java, PHP. The metric is the BLEU score. *We don't have Ruby in the pre-training corpus.

Model	All	Ruby	JavaScript	Go	Python	Java	PHP
CodeBERT [9]	17.83	12.16	14.90	18.07	19.06	17.65	25.16
PLBART [1]	18.32	14.11	15.56	18.91	19.30	18.45	23.58
ProphetNet-X [24]	18.54	14.37	16.60	18.43	17.87	19.39	24.57
CoTexT [22]	18.55	14.02	14.96	18.86	19.73	19.06	24.68
PolyglotCodeBERT [9]	19.06	14.75	15.80	18.77	18.71	18.71	20.11
DistillCodeT5 [39]	20.01	15.75	16.42	20.21	20.59	20.51	26.58
CodeGeeX (ours)	20.63	10.05*	16.01	24.62	22.50	19.60	31.00

model. Then, we test the performance of the fine-tuned model on each language pair with CodeBLEU score.

For all language pairs, we set the temperature to 0.2 and top-p to 0.95, and generate one translation for each sample in the test set. We report the results in Table 12. CodeGeeX performs better than all baseline models on all language pairs except for PHP to Python on the program level, C++ to Python on the snippet level, and PHP to Python on the snippet level. On average, CodeGeeX outperforms the baseline by 4.10 on the program level and by 1.99 on the snippet level.